Draft Doesn't Decide It: Champion Picks are Near-Uninformative for High-Elo League of Legends Win Prediction

A clean negative result for the draft, and a region- and patch-invariant signal from early objectives across NA, KR, and EUW

Yonghao Wang · DSC 148, University of California San Diego

Abstract

We ask what actually predicts the winner of a *high-elo* League of Legends game. We collect **51,000** apex-ladder matches (50,999 labeled) from three regions — NA, KR, EUW — via the Riot Match-V5 API (patches 16.9–16.11, Apr–May 2026), and predict whether the blue side wins. We compare two leakage-clean feature regimes: *draft-only* (champion picks, bans, summoner spells) and *draft + four first-objective* flags. **Finding 1 (negative):** at the top of the ladder the draft is near-uninformative — every draft-only model sits on the majority floor (AUC ≈ 0.54), and once features are matched, a gradient-boosted model beats logistic regression by only ΔAUC = +0.006 ($p = 0.17$). **Finding 2:** adding early objectives recovers a strong signal (AUC ≈ 0.79, first tower dominant) that is remarkably stable across regions (3×3 transfer gap 0.003) and patches ($p = 0.89$ for a patch feature). High-elo red side wins 54.2%, with KR closest to even — a descriptive MMR-matching effect. We report 1000× bootstrap confidence intervals throughout and a labeled leakage control (AUC 0.998). A live interactive demo is deployed online.

Keywords: esports analytics; win prediction; negative result; cross-region generalization; gradient boosting; SHAP; data leakage.

1. Introduction

The intuition behind League of Legends (LoL) win prediction is that the *draft* — which ten champions each team picks — should matter a great deal. Champions have lopsided matchups, and “draft diff” is a staple of community discourse. But is the draft actually informative once every player executes competently? We study this at the very top of the ranked ladder (Challenger/Grandmaster/Master), where mechanical execution is near-uniform and any predictive content of the draft should be most isolated.

We ask three questions. (i) How much of the outcome can the pre-game draft alone predict in high elo? (ii) How much does adding the earliest in-game objectives (first blood, tower, dragon, rift herald) recover? (iii) Does whatever signal exists generalize across regions and patches, or is it region-specific? To answer them cleanly we enforce a strict leakage discipline, match feature sets before comparing models, and attach 1000× bootstrap confidence intervals to every contrast so that tiny, sample-driven gaps are not mistaken for real effects.

Contributions. (1) A *clean negative result*: in high elo, champion picks are near-uninformative — draft-only models do not measurably beat the majority floor, and nonlinearity buys nothing once features are matched. (2) A *recovered, invariant signal*: early objectives lift AUC to ≈0.79 and that signal is nearly identical across three regions and three patches — what is predictable is region-agnostic game mechanics, while region differences live only in the (unpredictable) champion-preference meta. (3) A *descriptive* account of the high-elo red-side advantage and its regional spread. (4) A rigorous methodology — bootstrap significance, a labeled leakage control, feature-matched ablations — packaged with a deployed, interactive demo.

2. Related Work

Composition-only models. Prior work that predicts LoL outcomes from champion composition alone typically lands around 55–62% accuracy [2, 3]. A standard explanation is self-optimization: because players draft to counter and balance, the champion-vs-champion advantage is largely arbitraged away, pushing a pure-composition predictor toward the 50% coin-flip. Our high-elo setting is the extreme of this mechanism, and we indeed find even less signal than the 55–62% band.

Adding player skill. When models incorporate *player-on-champion mastery* (how much experience/performance a player has on the champion they locked), pre-game accuracy jumps to ≈75% [3, 4]; the player–champion term is repeatedly the single largest determinant. Combining pre-game and in-game state further improves over pre-game alone [4]. The lever we (deliberately) do not pull is player-on-champion mastery, which is not a direct Match-V5 endpoint; we return to it as the natural next step.

Side and region. A blue/red asymmetry is well documented, and at the top of the ladder the red-side edge is larger (community and Riot analyses put it near ~6% in Masters+), commonly attributed to MMR-matching dynamics and amplified by the unusually volatile 2026 apex ladder [1, 5]. We treat this as descriptive/correlational, not causal.

Our position. Relative to prior work we contribute a high-elo-specific study across three regions, a rigorously-tested *negative* result for the draft (with confidence intervals rather than point estimates), an explicit invariance test for the recovered objective signal, and a deployed demo for inspection.

3. Dataset and Exploratory Analysis

Collection. We seed from each region's apex ladder (Challenger–Master) and pull ranked matches through the Riot Match-V5 API with correct regional routing (NA→americas, KR→asia, EUW→europe). The corpus is **51,000** matches — 17,000 each from NA, KR, and EUW — on patches **16.9–16.11** (mostly 16.10), played **2026-04-29** to **2026-05-30**. We drop one match with *winner*=0, leaving **50,999** labeled games. Every *matchId* is round-tripped against the API and the final CSV contains **no player identifiers** (no PUUIDs or summoner names).

Leakage taxonomy. We partition fields into *safe* = draft (picks, bans, summoner spells); *allowed* = the four first-objective flags (blood, tower, dragon, rift herald); and *leaky* = end-of-game state (kill counts, first inhibitor, first baron, game duration). The leaky columns are permanently excluded from all real results (§6 quantifies why).

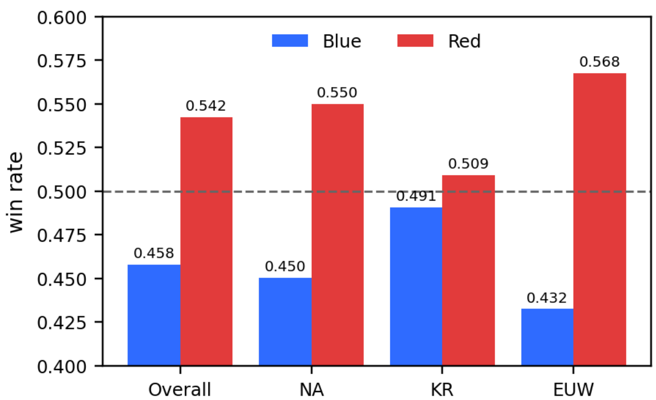

Figure 1. Win rate by side. High-elo red side wins **54.2%** overall (blue 45.8%). The bias is region-dependent — KR is closest to even (blue 49.1%), EUW most red-favored (blue 43.3%) — consistent with an MMR-matching account. *Design impact:* the majority floor is 54.2%, the bar every model must clear, and side is a useful prior, motivating distinct blue/red feature blocks.

(1) Side and region. Figure 1 shows red wins 54.2% overall (blue 45.8%), with a regional gradient: blue-win 49.1% (KR), 45.0% (NA), 43.3% (EUW). This sets the majority floor and motivates encoding side explicitly.

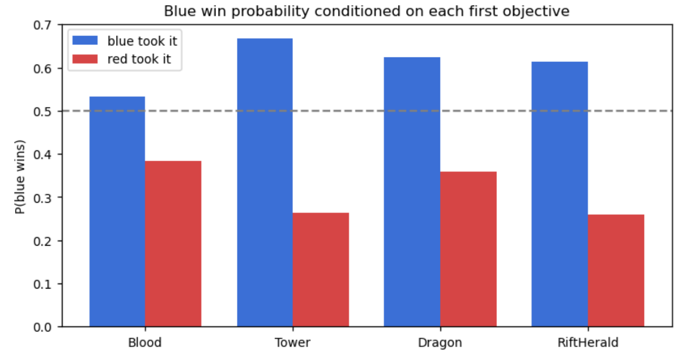

Figure 2. Blue-win probability conditioned on which side took each first objective. First tower is by far the strongest (gap **+0.405** between blue-took and red-took), then rift herald (+0.354), dragon (+0.265), blood (+0.148). *Design impact:* directly motivates the draft+objectives regime.

(2) Early objectives. Conditioning blue's win probability on who took each first objective (Fig. 2) reveals a strong, monotone signal: the blue–minus–red win gap is **+0.405** for first tower, +0.354 rift herald, +0.265 dragon, +0.148 first blood. This is the lever the draft lacks, and directly motivates regime B.

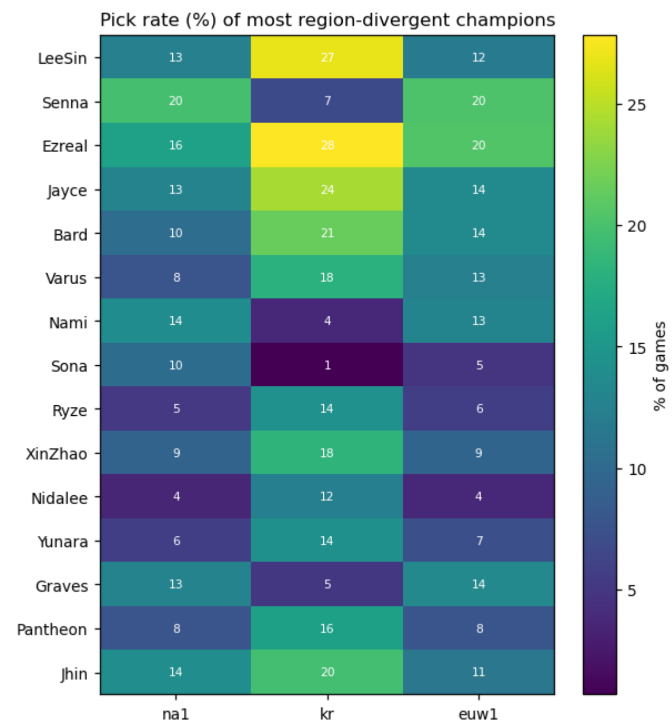

Figure 3. Pick rate (%) of the most region-divergent champions. KR is distinct (Lee Sin 27% vs 12–13%; Ezreal 28%; enchanters avoided — Sona 1%, Nami 4%). *Design impact:* region lives in champion *preferences*, which we show are exactly the unpredictable part.

(3) Regional meta divergence. Pick/ban rates differ sharply by region (Fig. 3); KR favors Lee Sin and avoids enchanters relative to NA/EUW. This frames our central story: regions differ in *what they draft*, yet §6 shows the draft is the part that does not predict outcomes.

4. Predictive Task

We predict a binary label: does blue (team 1 / teamId 100) win? We report accuracy, blue-class F1, and ROC-AUC; AUC is our primary metric because it is threshold- and prevalence-robust, which matters given the 54.2% majority floor that every model must beat. We define two feature **regimes**: **A (draft-only)** — multi-hot champion picks, bans, and summoner spells; and **B** —

regime A plus the four symmetric first-objective flags. The A→B contrast is the paper's main experiment.

Baselines. Logistic regression and Bernoulli Naive Bayes. They are apt foils for a composition signal: LR tests whether a *linear* combination of (binary) champion indicators separates wins, and NB tests a conditional-independence view of champion presence — both natural null models for "draft as a bag of champions." We hold the leaky columns out throughout, so no model can see post-hoc game state.

5. Model

Features. Champions are encoded as a *team-symmetric multi-hot* bag: within a team, pick order is irrelevant, but blue and red occupy separate column blocks so the model keeps a side prior. We never encode slot/lane positions — lane is a UI convenience, not a model input — which prevents memorizing positions and keeps the representation order-invariant. Region and patch enter as optional one-hot features.

Proposed models. A **LightGBM** gradient-boosted ensemble (optionally +region), and **ChampPoolNet**, a small network that learns a 32-d champion embedding and mean-pools it per team before a 64-unit head. The embedding-vs-one-hot choice is set up as a feature-matched ablation: the same head consumes either pooled embeddings or the raw multi-hot, isolating the representation. ChampPoolNet is far more parameter-efficient (≈149–157 input dims vs ≈709–717 for one-hot at equal AUC) — a promising compactness we flag as potential, not a measured accuracy gain on 172 champions.

Engineering. Three real difficulties shaped the pipeline. (i) LightGBM and PyTorch each bundle libomp; running them in one process *deadlocks* via duplicate OpenMP runtimes — fixed by single-threading and the duplicate-lib guard. (ii) Sparse vectorization of the multi-hot bag was needed to keep memory/time bounded. (iii) Under pandas 3.0 the new StringDtype turns empty cells into real NaN, which silently corrupted encodings until coerced. (iv) For the demo we found the all-objectives-*none* state is **0.000%** of training data — out-of-distribution for regime B — so the pre-game screen serves the genuinely-flat regime-A model and switches to B once any objective is set.

Optimization. LightGBM uses early stopping on a held-out split; a light sweep settled on learning rate 0.03 and 31 leaves (§6 shows insensitivity).

Application. We deploy the trained regime-B model behind a FastAPI service and a Next.js front-end: a user enters two drafts and toggles early objectives, and the blue win probability updates live with SHAP attributions. The pre-game screen is near-flat as champions are swapped (the negative result, made tangible); flipping *first tower* makes the number jump — the paper's thesis as an interaction.

6. Results

Model	A acc	A AUC	B acc	B F1	B AUC
LogReg	0.5423	0.5435	0.7215	0.6912	0.7846
BernoulliNB	0.5435	0.5448	0.7158	0.6940	0.7753
LightGBM	0.5504	0.5497	0.7232	0.6955	0.7914
ChampPoolNet	0.5429	0.5353	0.7236	0.6935	0.7884
OneHotMLP	0.5339	0.5439	0.7217	0.7058	0.7844

Table 1. Main results (all +region). Regime A (draft-only) vs B (draft+objectives). Majority floor (predict red) = 0.5422 acc. In regime A every AUC sits on the floor; in B all models reach ≈0.78–0.79. Numbers copied from *ckpt4_summary.txt*.

Finding 1 — the draft is near-uninformative (negative result). In regime A every model sits on the floor: AUC ranges 0.529–0.550 against a 0.5422 floor (Table 1). Crucially, once features are matched the nonlinear model does not help — LightGBM minus logistic regression is ΔAUC = –0.0017 (p = 0.668) without region and +0.0063 (p = 0.170) with it; learned embeddings minus one-hot is –0.0086 (p = 0.112) (Table 2). There is no detectable champion-interaction signal in high elo — even weaker than the 55–62% reported for lower elos, exactly as the self-optimization view predicts.

Contrast (ΔAUC)	Δ	95% CI	p
A: region alone (LGBM)	+.0092	[+.003,+.016]	.006*
A: nonlin–lin (LGBM–LR)	+.0063	[–.003,+.014]	.170
A: embed–onehot	–.0086	[–.019,+.002]	.112
B: LGBM–LR	+.0068	[+.004,+.010]	.000*
B: LGBM–NB	+.0161	[+.013,+.019]	.000*
B: embed–onehot	+.0040	[+.001,+.007]	.010*
B: +patch feature	+.0001	[–.001,+.001]	.892

Table 2. 1000× bootstrap paired AUC differences (* = 95% CI excludes 0). In regime A only the (tiny) region term is significant; nonlinearity and embeddings are within noise. In B the model gaps are significant but <0.02 AUC — statistically real, practically equivalent.

Finding 2 — objectives recover a strong, invariant signal. Regime B reaches AUC ≈ 0.79 (LightGBM 0.7914), squarely inside the 72–75% band reported for pre-game models — with first tower dominating the SHAP attributions. We emphasize the honest reading of Table 2: in regime B LightGBM is *statistically* the best (p<0.001 vs both baselines) but the margins are <0.02 AUC — the signal is in the *features, not the model*; the models are practically equivalent.

Region. Decomposed on AUC, region contributes a small but significant +0.0092 in regime A (p = 0.006) and essentially nothing in regime B (it is drowned out by the objectives). So the only predictive role of region is a weak side-rate prior, not a draft effect.

A (draft)	na1	kr	euw1	B (+obj)	na1	kr	euw1
na1	0.522	0.499	0.531	na1	0.769	0.794	0.785
kr	0.512	0.535	0.521	kr	0.766	0.798	0.782
euw1	0.502	0.505	0.524	euw1	0.766	0.784	0.781

Table 3. Cross-region 3×3 transfer (train-row × test-col, AUC). Regime A is indistinguishable from chance everywhere (diag 0.527 / off-diag 0.512). Regime B transfers almost perfectly: diagonal 0.783 vs off-diagonal 0.779, gap only **0.003**.

Cross-region transfer (Table 3). Using the strongest draft-only model, the A-matrix is flat at chance (mean diagonal 0.5271, off-diagonal 0.5116; the +0.0156 gap is within noise). The B-matrix is the headline invariance result: same-region 0.783 vs transfer 0.779, a 0.003 gap — the objective→win relationship is essentially identical across NA/KR/EUW. EUW-trained models lose accuracy on NA/KR (acc 0.658/0.641) while AUC holds (0.766/0.784): a calibration/threshold drift, so a deployment should re-calibrate its threshold per region rather than retrain.

Patch robustness. Per-patch AUC is 0.7888 / 0.7920 / 0.7963 for 16.9 / 16.10 / 16.11 (data is mostly 16.10); adding a patch one-hot moves AUC by +0.0001 (p = 0.892) — no effect. **Hyper-parameters.** Regime-B AUC stays within ±0.01 across learning rate {0.01–0.3}, leaves {15–127}, and estimators {100–3000}; the only thing that matters is using early stopping to avoid over-fitting (AUC drops to 0.781 at 3000 fixed trees).

Leakage control (not a real result). Adding the banned end-of-game columns to regime B inflates AUC from 0.7914 to **0.9983** (accuracy 0.721→0.981) — a +0.207 AUC jump. We report this only to show how post-hoc state would fabricate a near-perfect score, justifying the leakage discipline.

Case studies (regime B). (i) Most-confident-correct: NA1_5557961866, red took 3/4 objectives, P(blue)=0.056, red won. (ii) A confident upset: EUW1_7870164239, red took 3/4, P(blue)=0.077, yet *blue* won. (iii) The sharpest case — KR_8201503869: red swept **4/4** objectives, P(blue)=**0.098**, and blue still won. Together these show objectives are a strong but non-deterministic signal: top-ladder comebacks are real, an irreducible error floor that no pre/early feature set removes.

7. Conclusion

Three sentences summarize the study. In high elo the champion draft is near-uninformative — a clean negative result, consistent with and stronger than the 55–62% reported for lower elos. Early objectives recover a strong signal (AUC ≈ 0.79) that is invariant across three regions and three patches — what is predictable is region-agnostic game mechanics, while region differences live only in unpredictable champion preferences. The high-elo red-side advantage (54.2%, KR closest to even) is a descriptive MMR-matching pattern, not a causal claim.

Limitations. A single meta window (mostly patch 16.10) and, by design, no player-on-champion mastery feature. **Future work.** The literature's missing lever — Champion-Mastery and player-level history — is the natural next step; our negative draft result makes that lever the precise thing to add to push pre-game accuracy from ~0.54 toward the ~0.75 seen when player skill is modeled.

8. Availability and Reproducibility

Live demo. An interactive version is deployed at lol-win-predictor-web.onrender.com (a static Next.js front-end over a FastAPI service). Enter two drafts and toggle the early objectives: the blue win probability and its top SHAP factors update live. In the all-objectives-none state the number barely moves as champions are swapped — the negative result made tangible — and flipping *first tower* makes it jump from ≈46% to ≈77%.

Reproducibility. All results use a fixed seed (42) and a single stratified 80/20 split (50,999 → 40,799 train / 10,200 test), with the champion vocabulary fit on the training split only; every contrast in Table 2 is a 1000× bootstrap. The exact figures in Tables 1–3 are emitted verbatim by the checkpoint script (*models/ckpt4_summary.txt*); the trained models and serving code accompany this submission.

References

[1] Riot Games. *Riot Developer Portal* — Match-V5 API and regional routing (americas / asia / europe). developer.riotgames.com.

[2] Champion-composition win-prediction in MOBA games (~55–62% accuracy). Course-provided reference.

[3] Player-on-champion mastery as the dominant pre-game predictor (~75% accuracy). Course-provided reference.

[4] Combined pre-game and in-game outcome modeling for League of Legends. Course-provided reference.

[5] Blue/red side and regional win-rate asymmetry in high-elo League of Legends; MMR-matching dynamics. Course-provided reference.