

# Draft Doesn't Decide It: Champion Picks are Near-Uninformative for High-Elo League of Legends Win Prediction

A clean negative result for the draft, and a region- and patch-invariant signal from early objectives across NA, KR, and EUW

Yonghao Wang  
DSC 148, University of California San Diego  
La Jolla, CA, USA  
yonghaowang16@gmail.com

## ABSTRACT

We ask what actually predicts the winner of a *high-elo* League of Legends game. We collect **51,000** apex-ladder matches (50,999 labeled) from three regions—NA, KR, EUW—via the Riot Match-V5 API (patches 16.9–16.11, Apr–May 2026), and predict whether the blue side wins. We compare two leakage-clean feature regimes: *draft-only* (champion picks, bans, summoner spells) and *draft + four first-objective* flags. **Finding 1 (negative):** at the top of the ladder the draft is near-uninformative—every draft-only model sits on the majority floor ( $\text{AUC} \approx 0.54$ ), and once features are matched a gradient-boosted model does not beat logistic regression ( $\Delta\text{AUC} = -0.002$ ,  $p = 0.67$ ). **Finding 2:** adding early objectives recovers a strong signal ( $\text{AUC} \approx 0.79$ , first tower dominant) that is remarkably stable across regions ( $3 \times 3$  transfer gap 0.003) and patches ( $p = 0.89$  for a patch feature). High-elo red side wins 54.2%, with KR closest to even—a descriptive MMR-matching effect. We report 1000× bootstrap confidence intervals throughout and a labeled leakage control ( $\text{AUC} 0.998$ ). A live interactive demo is deployed at [lol-win-predictor-web.onrender.com](http://lol-win-predictor-web.onrender.com).

## KEYWORDS

esports analytics; win prediction; negative result; cross-region generalization; gradient boosting; SHAP; data leakage

## 1 INTRODUCTION

The intuition behind League of Legends (LoL) win prediction is that the *draft*—which ten champions each team picks—should matter a great deal. Champions have lopsided matchups, and “draft diff” is a staple of community discourse. But is the draft actually informative once every player executes competently? We study this at the very top of the ranked ladder (Challenger/Grandmaster/Master), where mechanical execution is near-uniform and any predictive content of the draft should be most isolated.

We ask three questions. (i) How much of the outcome can the pre-game draft alone predict in high elo? (ii) How much does adding the earliest in-game objectives (first blood, tower, dragon, rift herald) recover? (iii) Does whatever signal exists generalize across regions and patches, or is it region-specific? To answer them cleanly we enforce a strict leakage discipline, match feature sets before

comparing models, and attach 1000× bootstrap confidence intervals to every contrast so that tiny, sample-driven gaps are not mistaken for real effects.


**Contributions.** (1) A *clean negative result*: in high elo, champion picks are near-uninformative—draft-only models do not measurably beat the majority floor, and nonlinearity buys nothing once features are matched. (2) A *recovered, invariant signal*: early objectives lift AUC to  $\approx 0.79$  and that signal is nearly identical across three regions and three patches—what is predictable is region-agnostic game mechanics, while region differences live only in the (unpredictable) champion-preference meta. (3) A *descriptive* account of the high-elo red-side advantage and its regional spread. (4) A rigorous methodology—bootstrap significance, a labeled leakage control, feature-matched ablations—packaged with a deployed, interactive demo.

## 2 RELATED WORK

**Composition-only models.** Predicting LoL outcomes from champion composition *alone* is weak: community ML projects report  $\approx 55\%$  (Naive Bayes 55.3%) and  $\approx 53\%$  (tuned random forest) accuracy [2]. A standard explanation is self-optimization—because players draft to counter and balance, the champion-vs-champion advantage is largely arbitrated away, pushing a pure-composition predictor toward the 50% coin-flip. Our high-elo setting is the extreme of this mechanism.

**Adding player skill.** The signal lives in *player-on-champion mastery*, not the draft. Do et al. [3] reach **75.1%** accuracy at champion-select from players’ experience on their picked champions, concluding that individual champion skill matters “regardless of team composition”; summing each player’s per-champion win rate has been reported as high as  $\approx 93\%$  [3]. Peer-reviewed work that combines pre-game and in-game predictors improves further and notes early-game events are *less* predictive than late-game ones [4]. The lever we deliberately do not pull is player-on-champion mastery (not a direct Match-V5 endpoint); we return to it as the natural next step.

**Side and region.** A blue/red asymmetry is documented; community and industry analyses report a larger red edge at the top of the ladder—red  $\approx 53\%$  vs blue  $\approx 47\%$  in Master+ [5]. The proposed mechanism is MMR matching: at the apex ladder few players queue at once, so the two teams’ MMR cannot always be balanced, and the side seeded with the higher-MMR players (historically red)



**Figure 1: Win rate by side.** High-elo red side wins 54.2% overall (blue 45.8%). The bias is region-dependent—KR is closest to even (blue 49.1%), EUW most red-favored (blue 43.3%)—consistent with an MMR-matching account. *Design impact:* the majority floor is 54.2%, the bar every model must clear, and side is a useful prior, motivating distinct blue/red feature blocks.

wins more often [5]. This predicts a region with a larger, more evenly-matched apex pool should sit closest to 50/50—consistent with our KR observation (§3). We treat all of this as descriptive/correlation not causal.


**Our position.** Relative to prior work we contribute a high-elo-specific study across three regions, a rigorously-tested *negative* result for the draft (with confidence intervals rather than point estimates), an explicit invariance test for the recovered objective signal, and a deployed demo for inspection.

### 3 DATASET AND EXPLORATORY ANALYSIS


**Collection.** Seeded from each region’s apex league lists (Challenger–Master), we pull ranked matches through the Riot Match-V5 API with correct regional routing (NA→americas, KR→asia, EUW→europe), parallelized across the three regional hosts and deduplicated by matchId. The corpus is **51,000** matches—17,000 each from NA, KR, and EUW—on patches **16.9–16.11** (mostly 16.10), played **2026-04-29 to 2026-05-30**. We drop one match with *winner* = 0, leaving **50,999** labeled games. Every matchId is round-tripped against the API and the final CSV contains **no player identifiers** (no PUUIDs or summoner names).

**Leakage taxonomy.** We partition fields into *safe* = draft (picks, bans, summoner spells); *allowed* = the four first-objective flags (blood, tower, dragon, rift herald); and *leaky* = end-of-game state (kill counts, first inhibitor, first baron, game duration). The leaky columns are permanently excluded from all real results (§6 quantifies why).

**(1) Side and region.** Figure 1 shows red wins 54.2% overall (blue 45.8%), with a regional gradient: blue-win 49.1% (KR), 45.0% (NA), 43.3% (EUW). The ordering matches the apex MMR-matching account (§2): KR’s larger, more evenly-matched apex queue sits closest to 50/50, while thinner NA/EUW queues leave a bigger red edge. This sets the majority floor and motivates encoding side explicitly.



**Figure 2: Blue-win probability conditioned on which side took each first objective.** First tower is by far the strongest (gap +0.405 between blue-took and red-took), then rift herald (+0.354), dragon (+0.265), blood (+0.148). *Design impact:* directly motivates the draft+objectives regime.



**Figure 3: Pick rate (%) of the most region-divergent champions.** KR is distinct (Lee Sin 27% vs 12–13%; Ezreal 28%; enchanters avoided—Sona 1%, Nami 4%). *Design impact:* region lives in champion preferences, which we show are exactly the unpredictable part.

**(2) Early objectives.** Conditioning blue’s win probability on who took each first objective (Fig. 2) reveals a strong, monotone signal: the blue-minus-red win gap is +0.405 for first tower, +0.354 rift herald, +0.265 dragon, +0.148 first blood. This is the lever the draft lacks, and directly motivates regime B.

(3) **Regional meta divergence.** Pick/ban rates differ sharply by region (Fig. 3); KR favors Lee Sin and avoids enchanters relative to NA/EUW. This frames our central story: regions differ in *what they draft*, yet §6 shows the draft is the part that does not predict outcomes.

## 4 PREDICTIVE TASK

We predict a binary label: does blue (team 1 / teamId 100) win? We report accuracy, blue-class F1, and ROC-AUC; AUC is our primary metric because it is threshold- and prevalence-robust, which matters given the 54.2% majority floor that every model must beat. We define two feature **regimes**: **A (draft-only)**—multi-hot champion picks, bans, and summoner spells; and **B**—regime A plus the four symmetric first-objective flags. The A→B contrast is the paper's main experiment.

**Baselines.** Logistic regression and Bernoulli Naive Bayes. They are apt foils for a composition signal: LR tests whether a *linear* combination of (binary) champion indicators separates wins, and NB tests a conditional-independence view of champion presence—both natural null models for “draft as a bag of champions.” We hold the leaky columns out throughout, so no model can see post-hoc game state.

## 5 MODEL

**Features.** Champions are encoded as a *team-symmetric multi-hot* bag: within a team, pick order is irrelevant, but blue and red occupy separate column blocks so the model keeps a side prior. We never encode slot/lane positions—lane is a UI convenience, not a model input—which prevents memorizing positions and keeps the representation order-invariant. Region and patch enter as optional one-hot features.

**Proposed models.** A **LightGBM** gradient-boosted ensemble (optionally +region), and **ChampPoolNet**, a small network that learns a 32-d champion embedding and mean-pools it per team before a 64-unit head. The embedding-vs-one-hot choice is set up as a feature-matched ablation: the same head consumes either pooled embeddings or the raw multi-hot, isolating the representation. ChampPoolNet is far more parameter-efficient ( $\approx 149$ –157 input dims vs  $\approx 709$ –717 for one-hot at equal AUC)—a promising compactness we flag as potential, not a measured accuracy gain on 172 champions.

**Engineering.** Three real difficulties shaped the pipeline. (i) LightGBM and PyTorch each bundle libomp; running them in one process *deadlocks* via duplicate OpenMP runtimes—fixed by single-threading and the duplicate-lib guard. (ii) Sparse vectorization of the multi-hot bag was needed to keep memory/time bounded. (iii) Under pandas 3.0 the new StringDtype turns empty cells into real NaN, which silently corrupted encodings until coerced. (iv) For the demo we found the all-objectives-*none* state is **0.000%** of training data—out-of-distribution for regime B—so the pre-game screen serves the genuinely-flat regime-A model and switches to B once any objective is set.

**Optimization.** LightGBM uses early stopping on a held-out split; a light sweep settled on learning rate 0.03 and 31 leaves (§6 shows insensitivity).

**Application.** We deploy the trained regime-B model behind a FastAPI service and a Next.js front-end: a user enters two drafts

**Table 1: Main results (all +region). Regime A (draft-only) vs B (draft+objectives). Majority floor (predict red) = 0.5422 acc. In regime A every AUC sits on the floor; in B all models reach  $\approx 0.78$ –0.79. (OneHotMLP is the one-hot control arm of the embedding ablation; its slightly higher B-F1 is within noise and does not change the models' practical equivalence.) Numbers copied from ckpt4\_summary.txt.**

| Model        | A acc  | A AUC  | B acc  | B F1   | B AUC         |
|--------------|--------|--------|--------|--------|---------------|
| LogReg       | 0.5423 | 0.5435 | 0.7215 | 0.6912 | 0.7846        |
| BernoulliNB  | 0.5435 | 0.5448 | 0.7158 | 0.6940 | 0.7753        |
| LightGBM     | 0.5504 | 0.5497 | 0.7232 | 0.6955 | <b>0.7914</b> |
| ChampPoolNet | 0.5429 | 0.5353 | 0.7236 | 0.6935 | 0.7884        |
| OneHotMLP    | 0.5339 | 0.5439 | 0.7217 | 0.7058 | 0.7844        |

**Table 2: 1000× bootstrap paired AUC differences (\* = 95% CI excludes 0). In regime A only the (tiny) region term is significant; nonlinearity and embeddings are within noise. In B the model gaps are significant but <0.02 AUC—statistically real, practically equivalent.**

| Contrast ( $\Delta$ AUC) | $\Delta$ | 95% CI         | $p$   |
|--------------------------|----------|----------------|-------|
| A: region alone (LGBM)   | +0.092   | [+0.003,+0.16] | .006* |
| A: nonlin—lin (LGBM—LR)  | +0.063   | [−.003,+0.14]  | .170  |
| A: embed—onehot          | −.0086   | [−.019,+0.002] | .112  |
| B: LGBM—LR               | +0.068   | [+0.004,+0.10] | .000* |
| B: LGBM—NB               | +0.161   | [+0.13,+0.19]  | .000* |
| B: embed—onehot          | +0.040   | [+0.001,+0.07] | .010* |
| B: +patch feature        | +0.001   | [−.001,+0.001] | .892  |

and toggles early objectives, and the blue win probability updates live with SHAP attributions. The pre-game screen is near-flat as champions are swapped (the negative result, made tangible); flipping *first tower* makes the number jump—the paper's thesis as an interaction.

## 6 RESULTS

### Finding 1—the draft is near-uninformative (negative result).

In regime A every model sits on the floor: AUC ranges 0.529–0.550 against a 0.5422 floor (Table 1). Crucially, once features are matched the nonlinear model does not help—LightGBM minus logistic regression is  $\Delta$ AUC =  $-0.0017$  ( $p = 0.668$ )—this feature-matched, no-region contrast is the clean test of nonlinearity (the +0.0063,  $p = 0.170$ , with region folds in the region term); learned embeddings minus one-hot is  $-0.0086$  ( $p = 0.112$ ) (Table 2). There is no detectable champion-interaction signal in high elo. Measured as *lift over the majority-class baseline*, the gain is essentially nil—best regime-A accuracy is 0.5504 against the 0.5422 floor (under one point), and AUC 0.54–0.55 is barely above the 0.50 of random ranking—even smaller than the  $\approx +5$  points a (near-balanced) low-elo composition-only model gains over its 50% baseline [2]: self-optimization at its extreme.

**Finding 2—objectives recover a strong, invariant signal.** Regime B reaches AUC  $\approx 0.79$  (LightGBM 0.7914; accuracy 0.72), in line with the  $\approx 74\%$  accuracy reported for early-game / real-time models [4]—first tower dominates the SHAP attributions. We emphasize the honest reading of Table 2: in regime B LightGBM is *statistically* the best ( $p < 0.001$  vs both baselines) but the margins are

<0.02 AUC—the signal is in the *features*, not the *model*; the models are practically equivalent.

**Region.** Decomposed on AUC, region contributes a small but significant +0.0092 in regime A ( $p = 0.006$ ) and essentially nothing in regime B (it is drowned out by the objectives). So the only predictive role of region is a weak side-rate prior, not a draft effect.

**Cross-region transfer (Table 3).** Using the strongest draft-only model, the A-matrix is flat at chance (mean diagonal 0.5271, off-diagonal 0.5116; the +0.0156 gap is within noise). The B-matrix is the headline invariance result: same-region 0.783 vs transfer 0.779, a 0.003 gap—the objective→win relationship is essentially identical across NA/KR/EUW. EUW-trained models lose accuracy on NA/KR (acc 0.658/0.641) while AUC holds (0.766/0.784): a calibration/threshold drift, so a deployment should re-calibrate its threshold per region rather than retrain.

**Patch robustness.** Per-patch AUC is 0.7888 / 0.7920 / 0.7963 for 16.9 / 16.10 / 16.11 (data is mostly 16.10); adding a patch one-hot moves AUC by +0.0001 ( $p = 0.892$ )—no effect. **Hyper-parameters.** Regime-B AUC stays within  $\pm 0.01$  across learning rate {0.01–0.3}, leaves {15–127}, and estimators {100–3000}; the only thing that matters is using early stopping to avoid over-fitting (AUC drops to 0.781 at 3000 fixed trees).

**Leakage control (not a real result).** Adding the banned end-of-game columns to regime B inflates AUC from 0.7914 to **0.9983** (accuracy 0.721→0.981)—a +0.207 AUC jump. We report this only to show how post-hoc state would fabricate a near-perfect score, justifying the leakage discipline.

**Case studies (regime B).** (i) Most-confident-correct: NA1\_5557961866, red took 3/4 objectives,  $P(\text{blue})=0.056$ , red won. (ii) A confident upset: EUW1\_7870164239, red took 3/4,  $P(\text{blue})=0.077$ , yet *blue* won. (iii) The sharpest case—KR\_8201503869: red swept 4/4 objectives,  $P(\text{blue})=0.098$ , and blue still won. Together these show objectives are a strong but non-deterministic signal: top-ladder comebacks are real, an irreducible error floor that no pre/early feature set removes.


## 7 CONCLUSION

Three sentences summarize the study. In high elo the champion draft is near-uninformative—a clean negative result, consistent with and at the extreme of the  $\approx 55\%$  reported for composition-only models at lower elos [2]. Early objectives recover a strong signal (AUC  $\approx 0.79$ ) that is invariant across three regions and three patches—what is predictable is region-agnostic game mechanics, while region differences live only in unpredictable champion preferences. The high-elo red-side advantage (54.2%, KR closest to even) is a descriptive MMR-matching pattern, not a causal claim.

**Limitations.** A single meta window (mostly patch 16.10) and, by design, no player-on-champion mastery feature. **Future work.** The literature’s missing lever—Champion-Mastery and player-level history—is the natural next step; our negative draft result makes that lever the precise thing to add to push pre-game accuracy from  $\sim 0.54$  toward the  $\sim 0.75$  seen when player skill is modeled [3].

## 8 AVAILABILITY AND REPRODUCIBILITY

**Live demo.** An interactive version is deployed at `lol-win-predictor-web.onrender.com` (a static Next.js front-end over a FastAPI service). Enter two drafts



**Figure 4: Center readout of the deployed demo in the all-objectives-none state: regime A, 46.2% blue, “DRAFT IMPACT · LOW—champion composition barely moves the model.” The negative result as a live readout; setting *first tower* switches to regime B and the number jumps to  $\approx 77\%$ .**

and toggle the early objectives: the blue win probability and its top SHAP factors update live. In the all-objectives-none state the number barely moves as champions are swapped—the negative result made tangible (Fig. 4)—and flipping *first tower* makes it jump from  $\approx 46\%$  to  $\approx 77\%$ . (On the free tier the backend sleeps when idle; the first load may take  $\sim 30$ – $60$  s while it wakes.)

**Reproducibility.** All results use a fixed seed (42) and a single stratified 80/20 split (50,999  $\rightarrow$  40,799 train / 10,200 test), with the champion vocabulary fit on the training split only; every contrast in Table 2 is a 1000 $\times$  bootstrap. The exact figures in Tables 1–3 are emitted verbatim by the checkpoint script (`ckpt4_summary.txt`); the trained models and serving code accompany this submission.

## REFERENCES

- [1] Riot Games. *Riot Developer Portal—Match-V5 API and regional routing* (americas / asia / europe). `developer.riotgames.com`.
- [2] (community ML projects) J. Ou, *Smart Winner Predictor for League of Legends* (Weka,  $\sim 50k$  Riot-API matches): Naive Bayes 55.3% from team composition; J. Kang, *lol-ai* (GitHub): tuned random forest  $\approx 53\%$ .
- [3] T. D. Do, S. I. Wang, D. S. Yu, M. G. McMillian, and R. P. McMahan. *Using Machine Learning to Predict Game Outcomes Based on Player-Champion Experience in League of Legends*. Proc. 16th Int’l Conf. on the Foundations of Digital Games (FDG ’21), 2021. 75.1% at champion-select. arXiv:2108.02799. *Supporting*: T. Huang, D. Kim, V. Leung (2015), summed per-champion player win rates  $\approx 92.8\%$  (Naive Bayes).
- [4] *Applications of Linear and Ensemble-Based Machine Learning for Predicting Winning Teams in League of Legends*. Applied Sciences 15(10):5241, MDPI, 2025 (peer-reviewed): combining pre-game and in-game predictors improves; early-game events are less predictive than late. *Supporting: League of Legends: Real-Time Result Prediction*, arXiv:2309.02449, 2023 (LightGBM best,  $\approx 74.4\%$  on early-game data).
- [5] (community / industry analyses) *Red Side Advantage in High-Elo League of Legends*, Riftfeed, 2024—per League of Graphs, Master+ red  $\approx 53\%$  vs blue  $\approx 47\%$ , attributed to higher-MMR players placed on red; A. van Roon (Riot Games), official blue/red win rates by region, 2020 (via Esports Tales).

**Table 3: Cross-region  $3 \times 3$  transfer (train-row  $\times$  test-col, AUC). Regime A is indistinguishable from chance everywhere (diag 0.527 / off-diag 0.512). Regime B transfers almost perfectly: diagonal 0.783 vs off-diagonal 0.779, gap only 0.003. Numbers from ckpt4\_summary.txt.**

| A (draft-only) | na1   | kr    | euw1  | B (+objectives) | na1   | kr    | euw1  |
|----------------|-------|-------|-------|-----------------|-------|-------|-------|
| train na1      | 0.522 | 0.499 | 0.531 | train na1       | 0.769 | 0.794 | 0.785 |
| train kr       | 0.512 | 0.535 | 0.521 | train kr        | 0.766 | 0.798 | 0.782 |
| train euw1     | 0.502 | 0.505 | 0.524 | train euw1      | 0.766 | 0.784 | 0.781 |